transformers 4.41.2

Latest version

```
pip install transformers
```

Released: May 30, 2024

State-of-the-art Machine Learning for JAX, PyTorch and TensorFlow

Navigation

Project description

Release history

Download files

Verified details

These details have been verified by PyPI

Project description

English | 简体中文 | 繁體中文 | 한국어 | Español | 日本語 | हिन्दी | Русский | Português | తెలుగు | Français | Deutsch | Tiếng Việt |

State-of-the-art Machine Learning for JAX, PyTorch and TensorFlow

Maintainers


amysartran


ArthurZucker


lysandre


Thomwolf

Unverified details

*These details have **not** been verified by PyPI*

Project links

- Homepage

GitHub Statistics

- **Stars:** 127885
- **Forks:** 25364
- **Open issues:** 887
- **Open PRs:** 245

View statistics for this project via Libraries.io, or by using

🤗 Transformers provides thousands of pretrained models to perform tasks on different modalities such as text, vision, and audio.

These models can be applied on:

- 📝 Text, for tasks like text classification, information extraction, question answering, summarization, translation, and text generation, in over 100 languages.
- 🖼️ Images, for tasks like image classification, object detection, and segmentation.
- 🗣️ Audio, for tasks like speech recognition and audio classification.

Transformer models can also perform tasks on **several modalities combined**, such as table question answering, optical character recognition, information extraction from scanned documents, video classification, and visual question answering.

🤗 Transformers provides APIs to quickly download and use those pretrained models on a given text, fine-tune them on your own datasets and then share them with the community on our model hub. At the same time, each python module defining an architecture is fully standalone and can be modified to enable quick research experiments.

🤗 Transformers is backed by the three most popular deep learning libraries — Jax, PyTorch and TensorFlow — with a seamless integration between them. It's straightforward to train your models with one before loading them for inference with the other.

Online demos

our public dataset on Google BigQuery

Meta

License: Apache Software License (Apache 2.0 License)

Author: The Hugging Face team (past and future) with the help of all our contributors (https://github.com/huggingface/transformers/graphs/contributors)

NLP, vision, speech, deep, learning, transformer, pytorch, tensorflow, jax, BERT, GPT-2, Wav2Vec2, ViT

Requires: Python >=3.8.0

Classifiers

- **Development Status**
  - 5 - Production/Stable
- **Intended Audience**
  - Developers
  - Education
  - Science/Research
- **License**

You can test most of our models directly on their pages from the model hub. We also offer private model hosting, versioning, & an inference API for public and private models.

Here are a few examples:

In Natural Language Processing:

- Masked word completion with BERT
- Named Entity Recognition with Electra
- Text generation with Mistral
- Natural Language Inference with RoBERTa
- Summarization with BART
- Question answering with DistilBERT
- Translation with T5

In Computer Vision:

- Image classification with ViT
- Object Detection with DETR
- Semantic Segmentation with SegFormer
- Panoptic Segmentation with Mask2Former
- Depth Estimation with Depth Anything
- Video Classification with VideoMAE
- Universal Segmentation with OneFormer

In Audio:

- Automatic Speech Recognition with Whisper
- Keyword Spotting with Wav2Vec2
- Audio Classification with Audio Spectrogram Transformer

- OSI Approved :: Apache Software License
- **Operating System**
  - OS Independent
- **Programming Language**
  - Python :: 3
  - Python :: 3.8
  - Python :: 3.9
  - Python :: 3.10
- **Topic**
  - Scientific/Engineering :: Artificial Intelligence

In Multimodal tasks:

- Table Question Answering with TAPAS
- Visual Question Answering with ViLT
- Image captioning with LLaVa
- Zero-shot Image Classification with SigLIP
- Document Question Answering with LayoutLM
- Zero-shot Video Classification with X-CLIP
- Zero-shot Object Detection with OWLv2
- Zero-shot Image Segmentation with CLIPSeg
- Automatic Mask Generation with SAM

100 projects using Transformers

Transformers is more than a toolkit to use pretrained models: it's a community of projects built around it and the Hugging Face Hub. We want Transformers to enable developers, researchers, students, professors, engineers, and anyone else to build their dream projects.

In order to celebrate the 100,000 stars of transformers, we have decided to put the spotlight on the community, and we have created the awesome-transformers page which lists 100 incredible projects built in the vicinity of transformers.

If you own or use a project that you believe should be part of the list, please open a PR to add it!

If you are looking for custom support from the Hugging Face team